



# Understanding the problem(s) Fabric solves

The story of Houston Electrics Inc



# Understanding the problem(s) Fabric solves

## The story of Houston Electrics Inc



Based in the USA, ships electrical parts around the world

400+ employees

Innovative and data-driven

- 2015: Digital transformation programme to drive growth
  - Central data department with 50 employees:
  - Investment in Azure, AWS and Power BI for reporting

## Existing tools @ Houston Electric

Customer Success Team



Data Engineering Team



IT Team

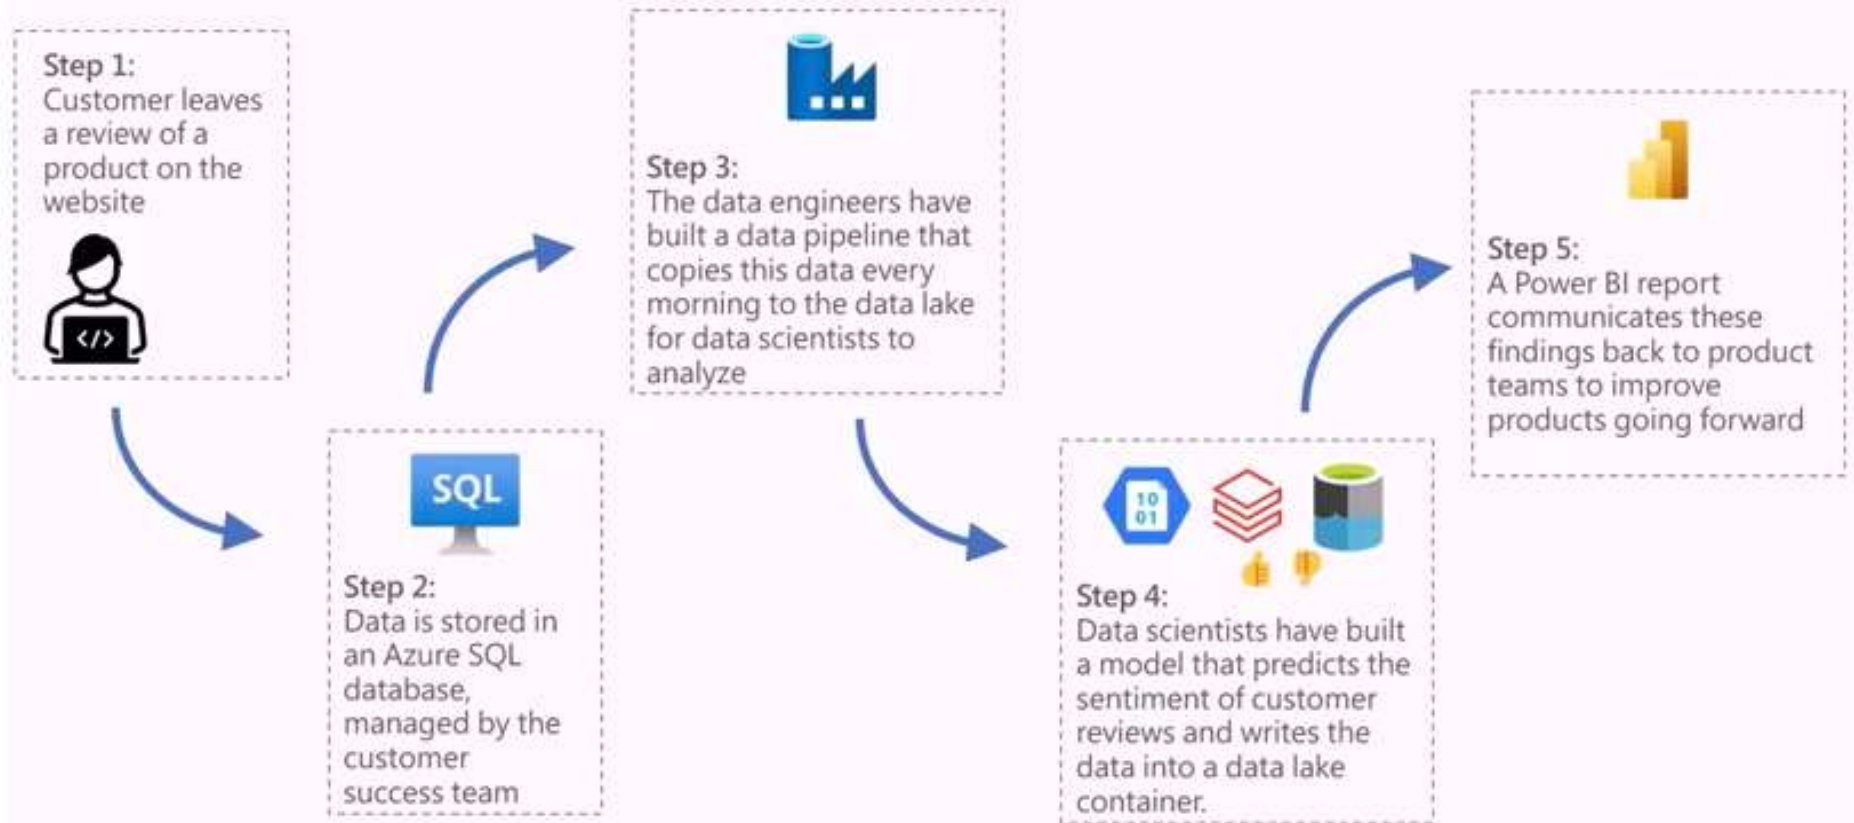
Data Science Team



Manufacturing Team



## Existing customer reviews workflow



## Existing tools @ Houston Electric



Customer Success Team



100010111

100010111

100010111

Data Engineering Team



100010111



IT Team

100010111

Data Science Team



100010111

Manufacturing Team

100010111  
100010111







## IF YOU SPEAK TO...



 Chief Data Officer

"

*The architecture has grown organically and now there are data silos all over the place.*

*I am the Chief Data Officer and don't want to be the Chief Integration Officer*

"





## IF YOU SPEAK TO...



Data Engineering Lead:

"

*We're maintaining hundreds of pipelines  
copying data between lots of different  
data stores for different departments.  
It's messy*

"



## IF YOU SPEAK TO...

 Data Scientist:

"

*We have data scattered in so many formats in different places, it takes me days just to get clean datasets just to begin an analysis. Even then, I don't know if I can trust the data.*

"

## IF YOU SPEAK TO...

 Data Scientist:

"


*We have data scattered in so many formats in different places, it takes me days just to get clean datasets just to begin an analysis. Even then, I don't know if I can trust the data.*

"  
"

*I had to learn the intricacies of many different data technologies, and each one is different. We had a new starter. It took 3 months to upskill them in all the different platforms we use.*

"





*"We're using too many systems, all have different security profiles and requirements to keep data at-rest and in-transit secure. It's a nightmare"*





## IF YOU SPEAK TO...



IT Director

"

*I dread getting our Azure bill every month. So unpredictable, and sometimes scary. Each data product has their own pricing structure so it's difficult to predict how much we will be charged month-to-month.*

"

## IF YOU SPEAK TO...

 Dashboard user:

"

*I'm not sure if I can trust the data  
that is being presented to me –  
doesn't always reflect reality*

"

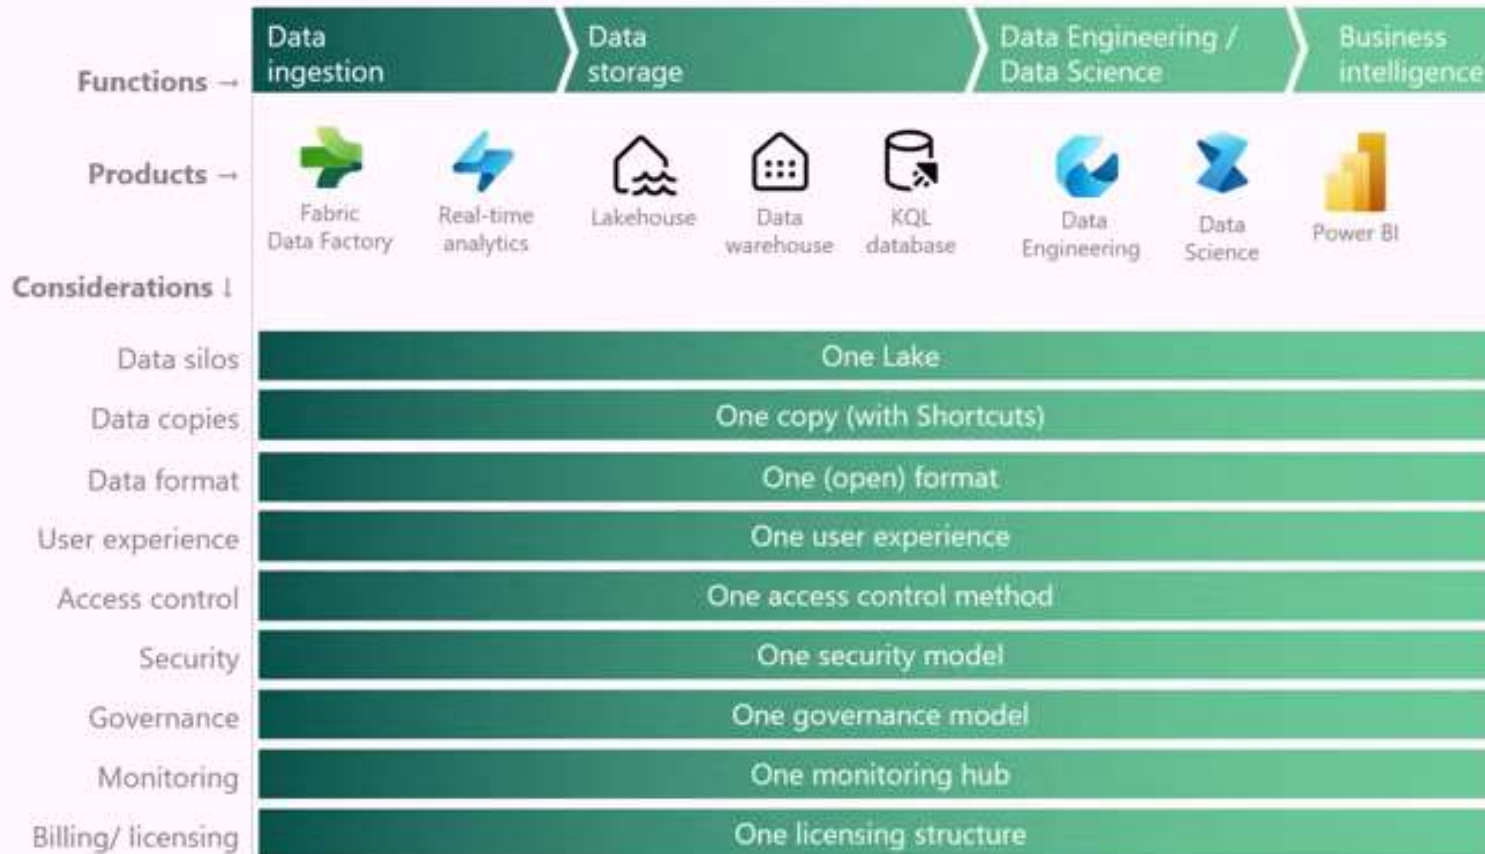
# HOW FABRIC IS DIFFERENT



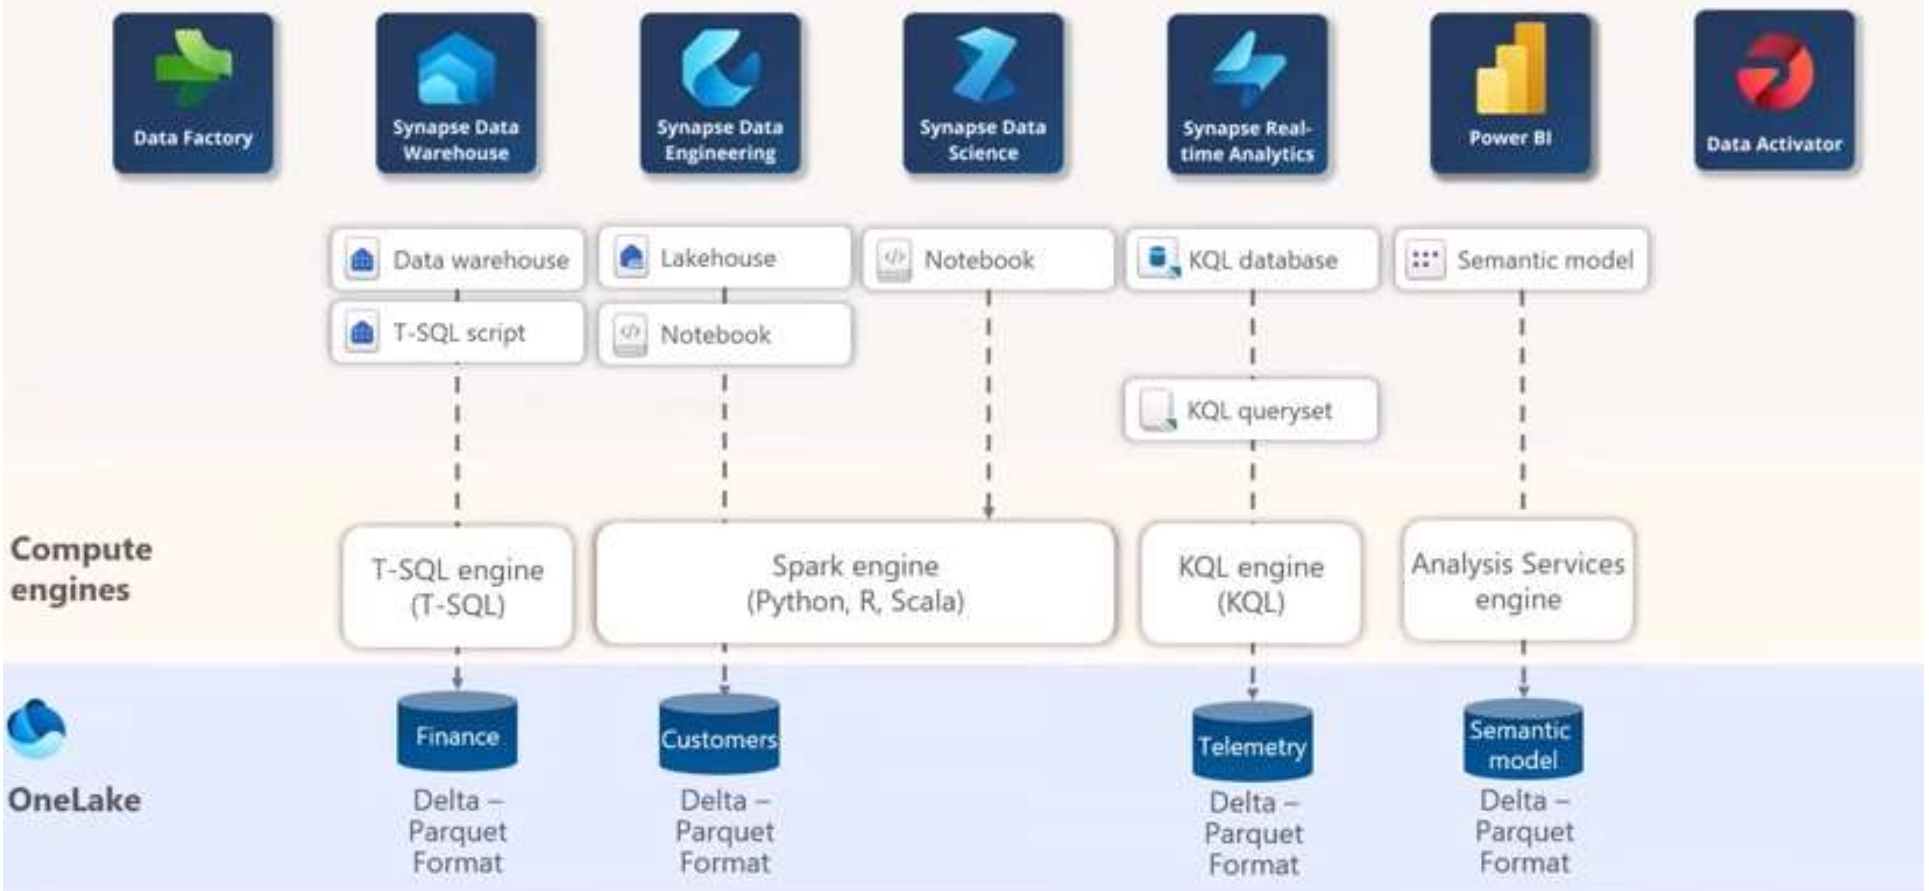
# HOW FABRIC IS DIFFERENT



# HOW FABRIC IS DIFFERENT



## Fabric experiences





# DATA FACTORY

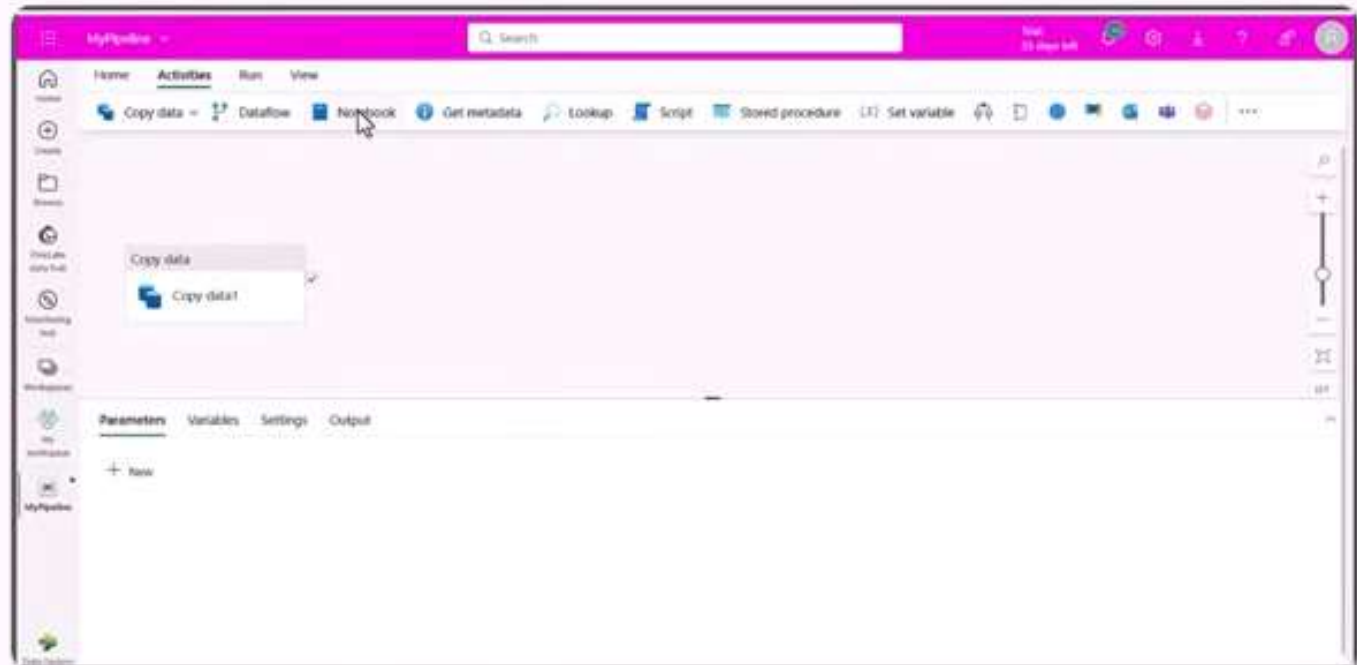
## Core purpose:

Moving and transforming your data. A set of tools to help you with Extract, Transform and Load.

## Fabric items:

 Dataflow

 Data pipeline



# DATA FACTORY

## **Core purpose:**

Moving and transforming your data. A set of tools to help you with Extract, Transform and Load.

## **Fabric items:**



Dataflow



Data pipeline

## **Similar to:**

- Azure Data Factory
- Synapse Pipelines
- Power BI Dataflow Gen1

## **Main personas:**

- Data engineers
- Analytics Engineers
- Power BI developers



# SYNAPSE DATA WAREHOUSE

## Core purpose:

provides a familiar transactional data warehouse solution with tables, schemas, views, stored procedures etc. Query-able using Structured Query Language (T-SQL)

The screenshot displays the Synapse Data Warehouse user interface. On the left is a navigation pane with a tree view showing the hierarchy: Home, Explorer, Warehouses, Schemas, dbo, Tables, Date, Geography, HackneyLicense, Medication, Time, Trip, Weather, Views, Functions, and Stored Procedures. The main workspace is titled 'Visual query 1' and shows a diagram of a 'Trip' table with a 'Source' column. A context menu is open over the 'Source' column, offering options: 'Keep top rows', 'Keep bottom', 'Keep range of rows', 'Keep duplicates', 'Keep errors', and 'Filter rows'. Below the diagram, there are buttons for 'Save as table', 'Download Excel file', and 'Visualize results'. At the bottom, a table displays query results with columns: TripID, MedicationID, HackneyLicenseID, PickupTimeID, DropoffTimeID, PickupGeographyID, DropoffGeographyID, and PickUpLatitude. The table contains 7 rows of data.

TripID	MedicationID	HackneyLicenseID	PickupTimeID	DropoffTimeID	PickupGeographyID	DropoffGeographyID	PickUpLatitude
1	20130021	0007	27470	0809	20460	8408	40.7477
2	20130025	0206	3622	20275	21008	26032	40.7204
3	20130400	1108	2088	16123	8090	14081	40.7088
4	20131008	4137	23101	40440	34000	10084	40.7631
5	20130810	7189	13108	4889	62062	7740	40.7452
6	20130000	9123	40188	20300	21080	22125	40.7321
7	20130408	8413	28764	17116	16021	22712	40.7087



# SYNAPSE DATA WAREHOUSE

**Core purpose:**

provides a familiar transactional data warehouse solution with tables, schemas, views, stored procedures etc. Query-able using Structured Query Language (T-SQL)

**Fabric items:**

Data warehouse

**Similar to:**

- SQL Server/ Azure SQL
- Synapse SQL Serverless/Dedicated
- Snowflake

**Main personas:**

- Database administrators
- Data engineers
- Data analysts

# SYNAPSE DATA ENGINEERING

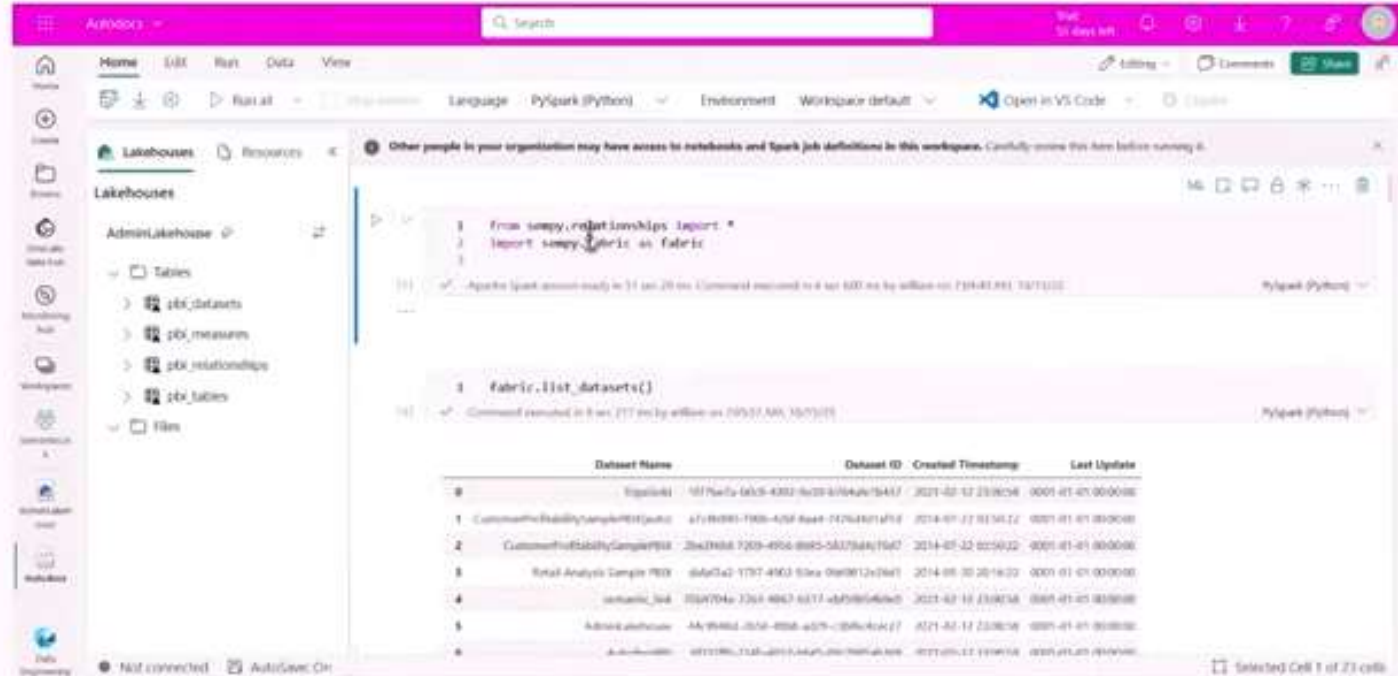
## Core purpose:

Enables users to design, build, and maintain **infrastructures and systems** that enable their organizations to collect, store, process, and analyze large volumes of data .

## Fabric items:

 Lakehouse

 Notebook



The screenshot displays the Synapse Data Engineering interface. On the left, a sidebar shows the 'Lakehouses' section with a tree view containing 'AdminLakehouse' and 'Tables'. The 'Tables' section is expanded, showing a list of tables: 'tbl\_datasets', 'tbl\_measures', 'tbl\_relationships', and 'tbl\_tables'. The main area shows a notebook with two code cells. The first cell contains the following code:

```
1 from synapse.relationships import *
2 import synapse.fabric as fabric
3
```

The second cell contains the following code:

```
1 fabric.list_datasets()
```

Below the code cells, a table of datasets is displayed. The table has four columns: 'Dataset Name', 'Dataset ID', 'Created Timestamp', and 'Last Update'. The table contains six rows of data.

	Dataset Name	Dataset ID	Created Timestamp	Last Update
0	tbl_datasets	tbl_datasets-4000-4000-4000-4000	2021-02-12 23:00:00	2021-02-12 23:00:00
1	CustomerRelationshipSamplePilot	tbl_datasets-4000-4000-4000-4000	2021-02-12 23:00:00	2021-02-12 23:00:00
2	CustomerRelationshipSamplePilot	tbl_datasets-4000-4000-4000-4000	2021-02-12 23:00:00	2021-02-12 23:00:00
3	CustomerRelationshipSamplePilot	tbl_datasets-4000-4000-4000-4000	2021-02-12 23:00:00	2021-02-12 23:00:00
4	CustomerRelationshipSamplePilot	tbl_datasets-4000-4000-4000-4000	2021-02-12 23:00:00	2021-02-12 23:00:00
5	CustomerRelationshipSamplePilot	tbl_datasets-4000-4000-4000-4000	2021-02-12 23:00:00	2021-02-12 23:00:00

# SYNAPSE DATA ENGINEERING

## Core purpose:

Enables users to design, build, and maintain **infrastructures and systems** that enable their organizations to collect, store, process, and analyze large volumes of data .

## Fabric items:

-  Lakehouse
-  Notebook
-  Spark Job

## Similar to:

- Azure Data Lake Storage\* (ADLS Gen2)
- Databricks
- Snowflake

## Main personas:

- Data engineers
- Analytics Engineers

# SYNAPSE DATA SCIENCE

## Core purpose:

Supports the entire data science workflow in an organisation, from data exploration, preparation and cleansing to experimentation, modeling, model scoring and serving of predictive insights to BI reports.

## Fabric items:

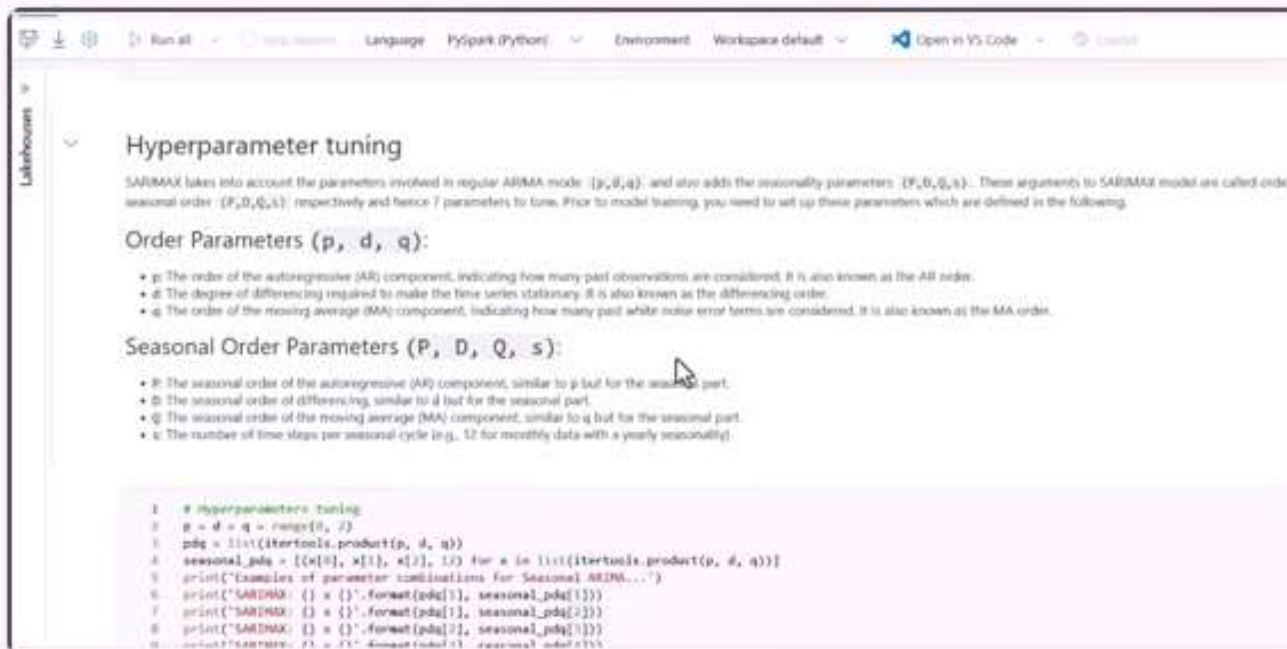
 Notebook

 Experiments

 ML Models

## Similar to:

- Azure Machine Learning
- Synapse notebooks
- Databricks notebooks



**Hyperparameter tuning**

SARIMAX takes into account the parameters involved in regular ARIMA model:  $(p, d, q)$ , and also adds the seasonality parameters:  $(P, D, Q, s)$ . These arguments to SARIMAX model are called order seasonal order:  $(P, D, Q, s)$  respectively and hence 7 parameters to tune. Prior to model training, you need to set up these parameters which are defined in the following.

**Order Parameters  $(p, d, q)$ :**

- $p$ : The order of the autoregressive (AR) component, indicating how many past observations are considered. It is also known as the AR order.
- $d$ : The degree of differencing required to make the time series stationary. It is also known as the differencing order.
- $q$ : The order of the moving average (MA) component, indicating how many past white noise error terms are considered. It is also known as the MA order.

**Seasonal Order Parameters  $(P, D, Q, s)$ :**

- $P$ : The seasonal order of the autoregressive (AR) component, similar to  $p$  but for the seasonal part.
- $D$ : The seasonal order of differencing, similar to  $d$  but for the seasonal part.
- $Q$ : The seasonal order of the moving average (MA) component, similar to  $q$  but for the seasonal part.
- $s$ : The number of time steps per seasonal cycle (e.g., 12 for monthly data with a yearly seasonality).

```
1 # Hyperparameters tuning
2 p = d = q = range(0, 2)
3 pdq = list(itertools.product(p, d, q))
4 seasonal_pdq = [(x[0], x[1], x[2], 12) for x in list(itertools.product(p, d, q))]
5 print("Examples of parameter combinations for Seasonal ARIMA...")
6 print("SARIMAX: {} = {}".format(pdq[1], seasonal_pdq[1]))
7 print("SARIMAX: {} = {}".format(pdq[2], seasonal_pdq[2]))
8 print("SARIMAX: {} = {}".format(pdq[3], seasonal_pdq[3]))
9 print("SARIMAX: {} = {}".format(pdq[4], seasonal_pdq[4]))
```



## SYNAPSE REAL-TIME ANALYTICS


### Core purpose:

Provides a set of tools to ingest, manage and analyze real-time event data

### Fabric items:

 KQL database

 Eventstream

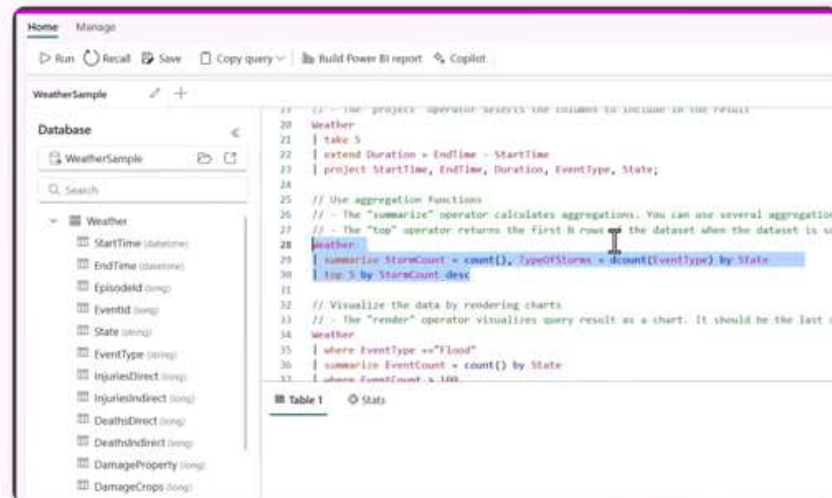
 KQL queryset

### Similar to:

- Azure Data Explorer

### Main personas:

- Data engineers  
- Analytics Engineers



# POWER BI

## Core purpose:

Power BI is Microsoft's business intelligence solution that allows you to create reports to present visual insights to business users.

## Fabric items:



Report



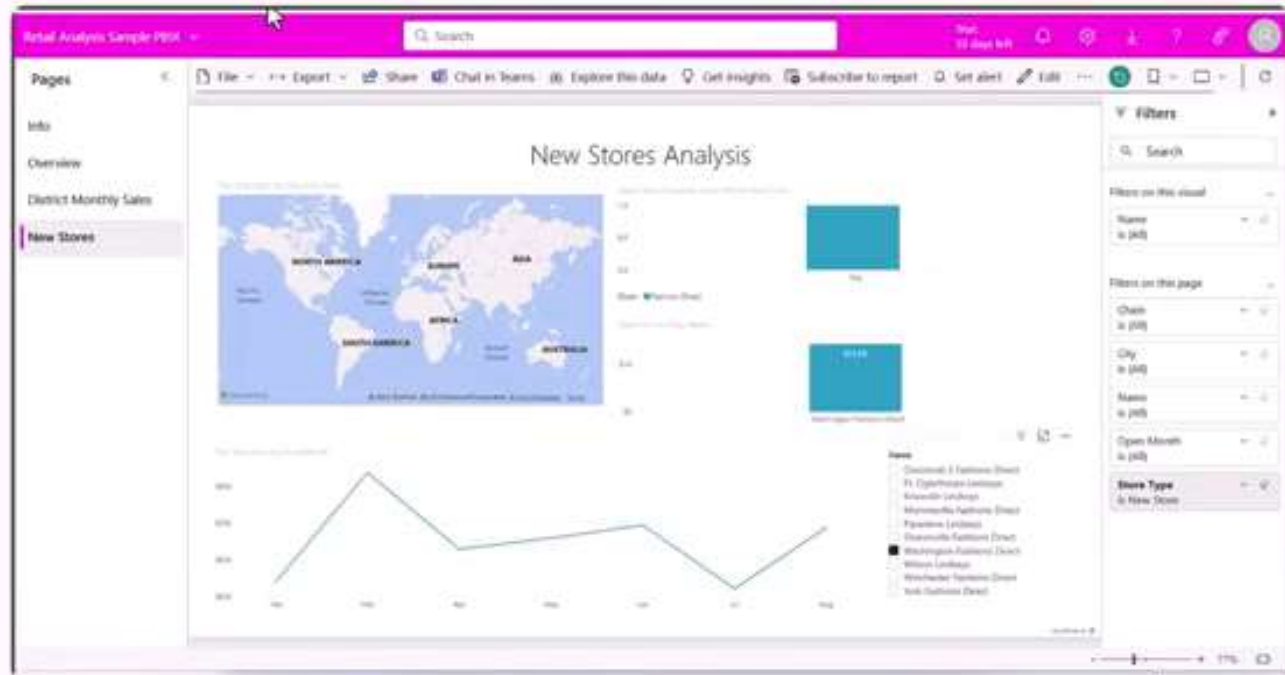
Semantic model

## Similar to:

- Tableau
- Looker

## Main personas:

- Business users
- Power BI developers
- Data analysts



# DATA ACTIVATOR

## Core purpose:

Automatically taking actions (like running a Power Automate routine) when **patterns** or **conditions** are detected in changing data, such as data in Power BI reports and Eventstreams

## Fabric items:



## Similar to:

- Power Automate
- Azure Functions

## Main personas:

- Business users
- Power BI developers
- Data analysts

